



TensorFlow

SPEAKER: PYRY TAKALA

True AI

DIALOGUES +
DEEP LEARNING

WAITING TIMES



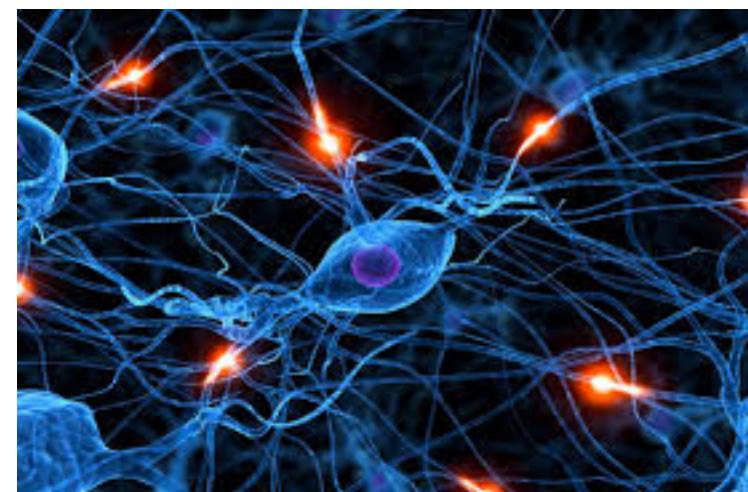
“...it may take up to 20 working days to provide a full reply where investigation is required. During times of high demand, we may sometimes miss our targets...”

- London Midland trains

LONDON-BASED STARTUP

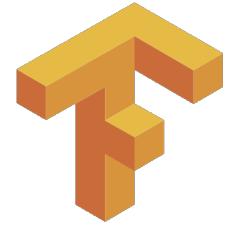


DEEP LEARNING RESEARCHERS



EXPERT TEAM



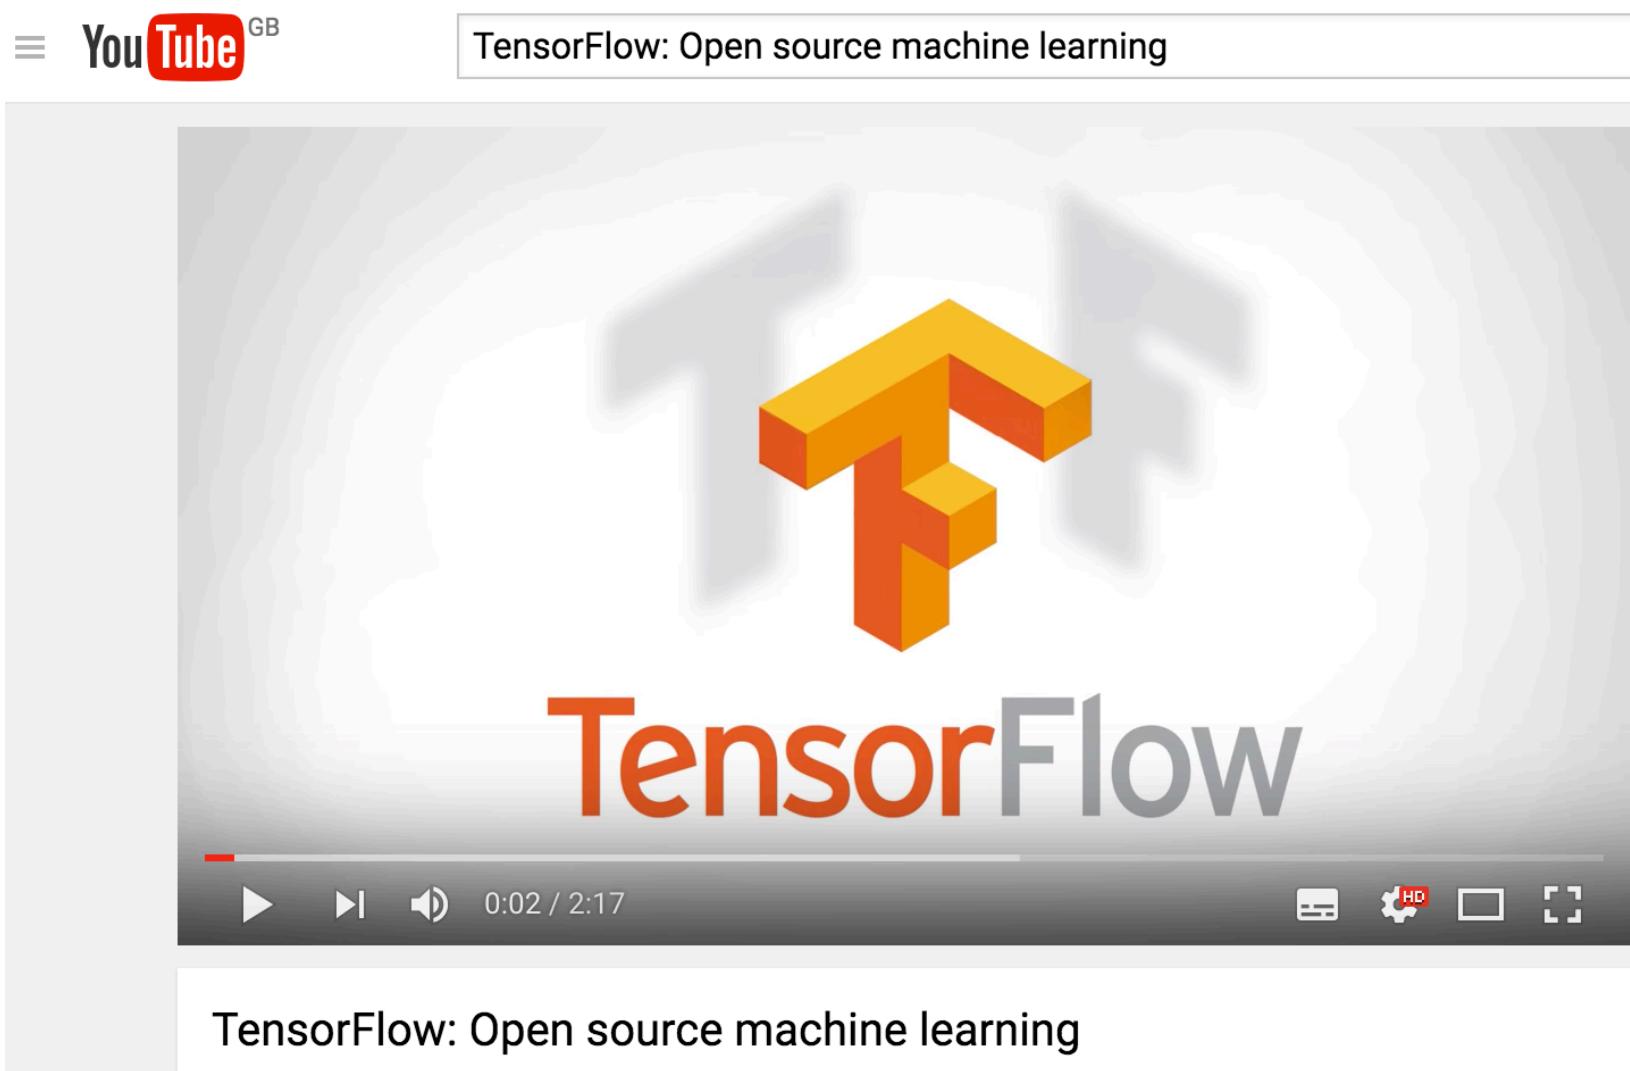


Background

Usage

Under-the-hood

Resources

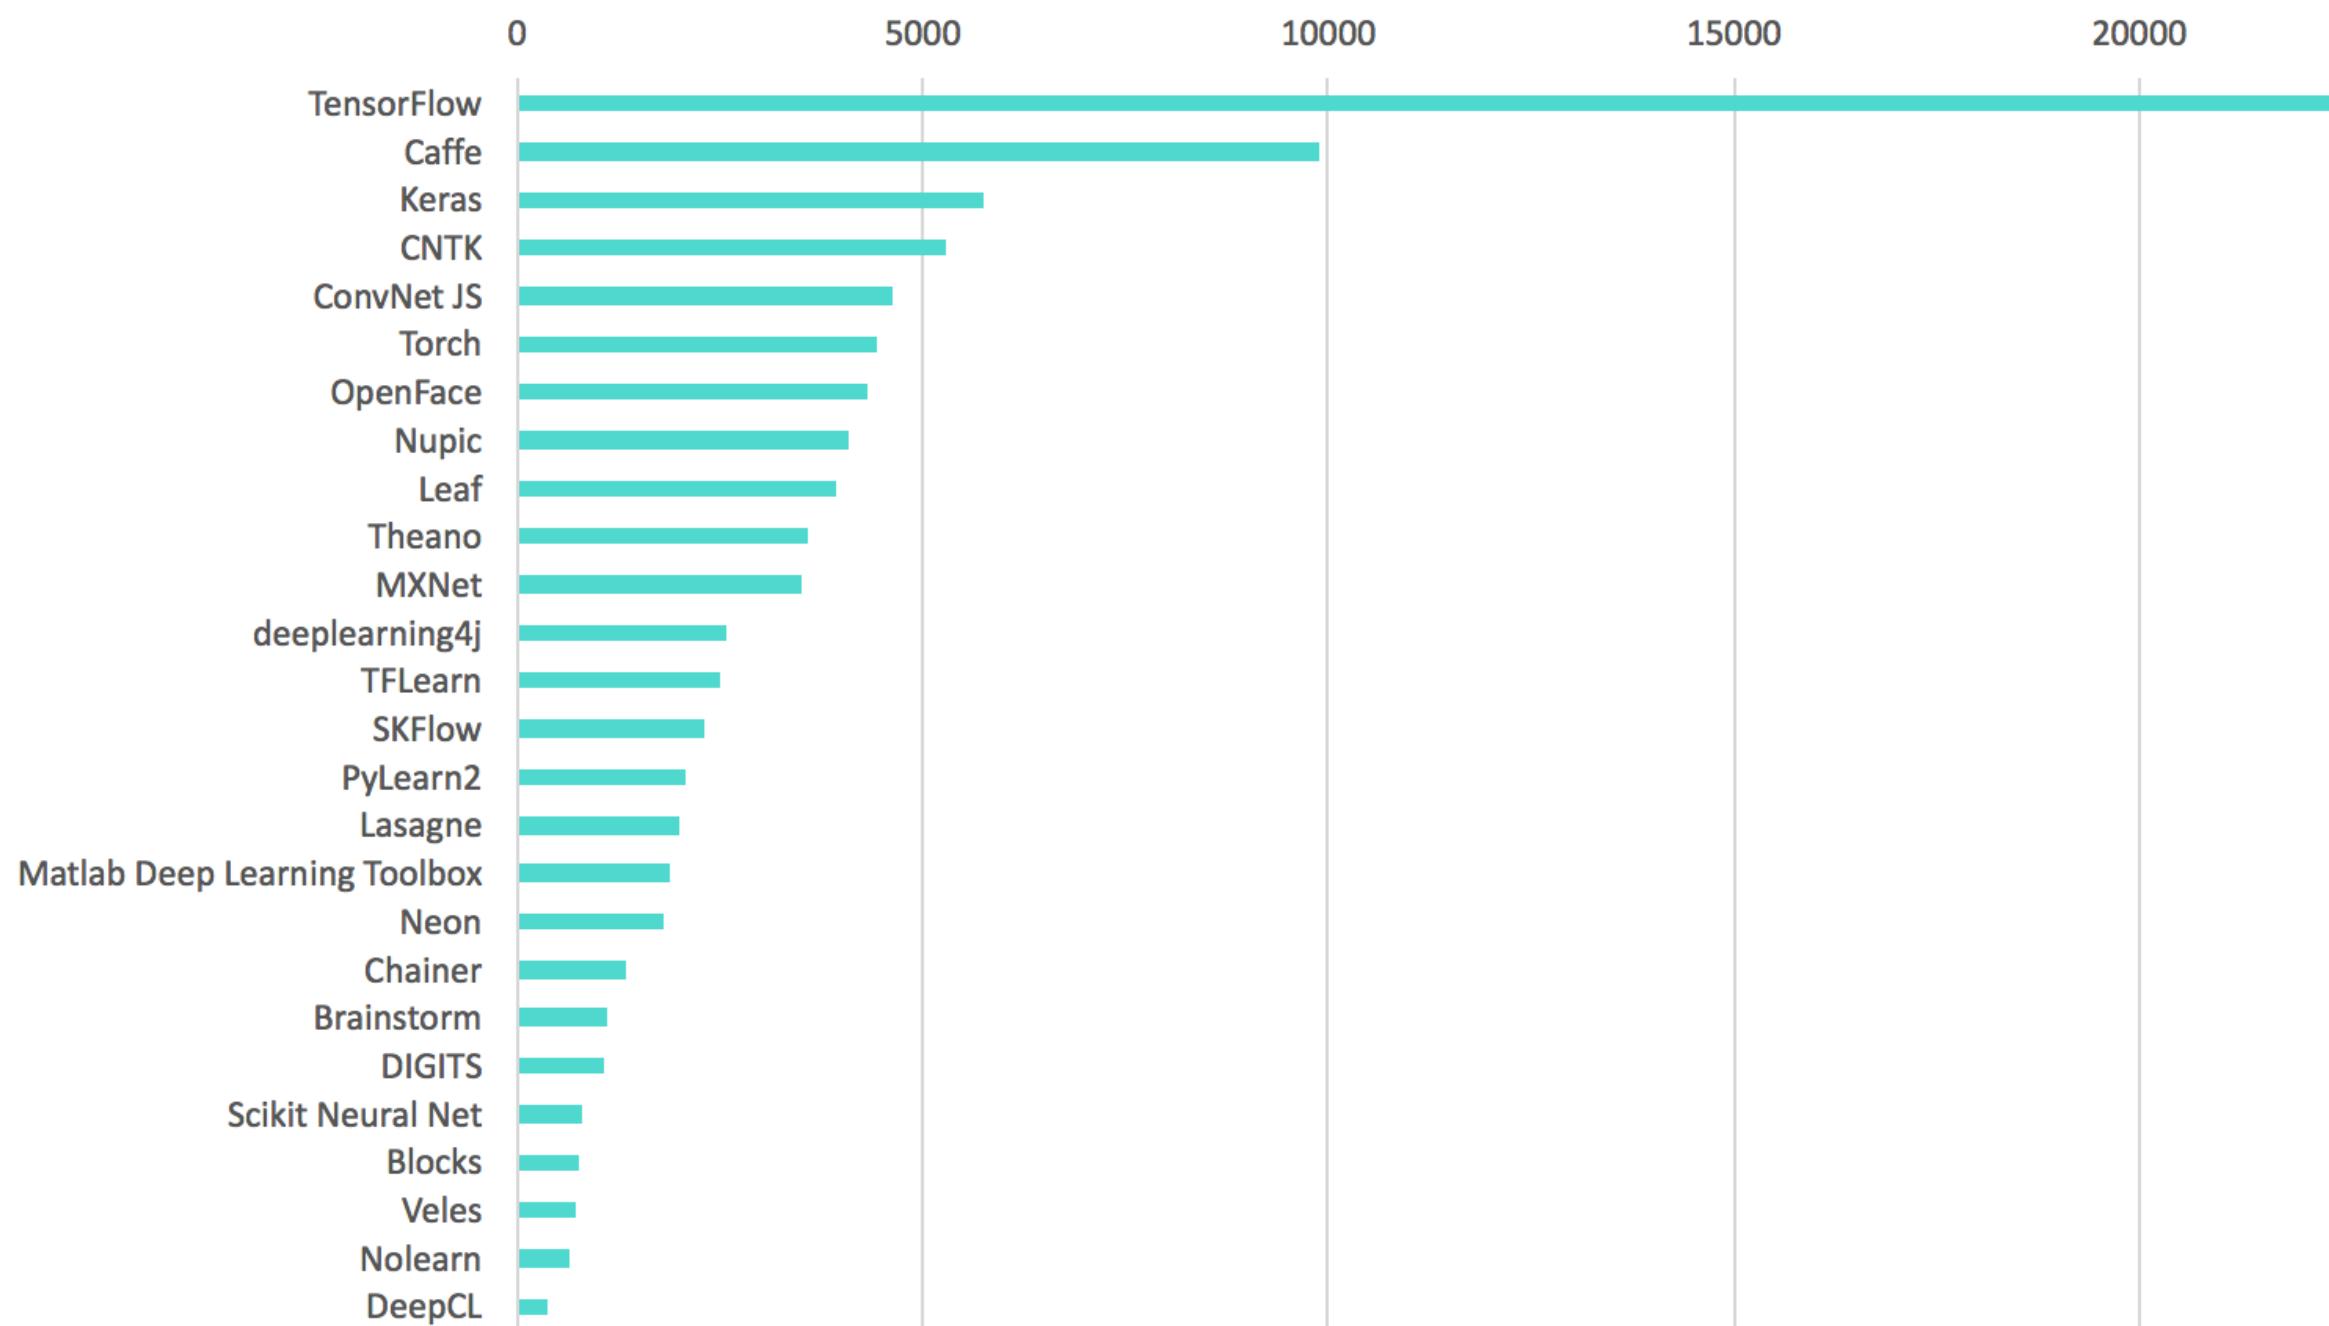


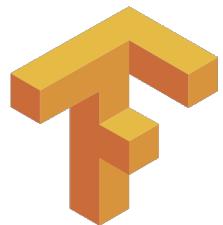
TensorFlow: Open source machine learning

https://www.youtube.com/watch?v=oZikw5k_2FM

- ✓ Research and production ML library, designed especially for deep learning
- ✓ Started by Google Brain after DistBelief; now open source
- ✓ Strong focus on parallelizing computations
- ✓ Python & C++

GITHUB STARS PER FRAMEWORK





Pros and Cons

- (+) Python + numpy
- (+) Computational graph abstraction, like Theano
- (+) Much faster compile times than Theano
- (+) TensorBoard for visualization
- (+) Data AND model parallelism
- (-) Slower than other frameworks
- (-) Much "fatter" than Torch; more magic
- (-) Not many pretrained models
- (-) Computational graph is pure Python, so slow

theano

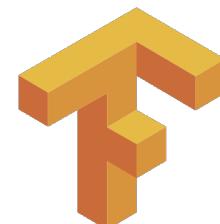
Pros and Cons

- (+) Python + numpy
- (+) Computational graph is nice abstraction
- (+) RNNs fit nicely in computational graph
- (-) Raw Theano is somewhat low-level
- (+) High level wrappers (Keras, Lasagne) ease the pain
- (-) Error messages can be unhelpful
- (-) Large models can have long compile times
- (-) Much "fatter" than Torch; more magic
- (-) Patchy support for pretrained mo



Pros and Cons:

- (-) Lua
- You usually write your own training code (Less plug and play)
- (+) Lots of modular pieces that are easy to combine
- (+) Easy to write your own layer types and run on GPU
- (+) Lua. ;) (Most of the library code is in Lua, easy to read)
- (+) Lots of pretrained models!
- (-) Not good for recurrent neural networks



theano



Modeling capability	***	**	****'	****'	*****
Interfaces	***	**'	****'	***	***
Model Deployment	*****	****'	****'	***	***
Performance			*****	***	*****
Architecture	***		*****	***	*****
Ecosystem	C++	C++	Py & C++	Py	Lua
Cross-platform					

WHAT IS A “TENSOR”:

“Tensors are geometric objects that describe linear relations between geometric vectors, scalars, and other tensors.”

W

A scalar is a tensor

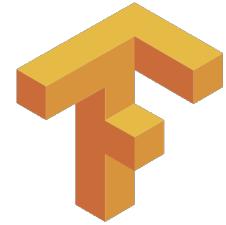
24

A vector is a tensor

[24 98 16]

A matrix is a tensor

$\begin{bmatrix} 24 & 98 & 16 \\ 5 & 21 & 2 \end{bmatrix}$



Background

Usage

Under-the-hood

Resources

```
In [1]: import tensorflow as tf
```

```
In [2]: with tf.Session() as sess:  
....:     print tf.constant("Hello world!").eval()  
....:
```

```
Hello world!
```

NUMPY VS TENSORFLOW

```
In [3]: import numpy as np  
  
In [4]: a = np.zeros((2,2))  
  
In [5]: b = np.ones((2,2))  
  
In [6]: print np.sum(b, axis=1)  
[ 2.  2.]
```

```
In [1]: import tensorflow as tf  
  
In [2]: with tf.Session() as sess:  
....:     a = tf.zeros((2,2))  
....:     b = tf.ones((2,2))  
....:     print tf.reduce_sum(b, reduction_indices=1).eval()  
....:  
[ 2.  2.]
```

NUMPY VS TENSORFLOW

Numpy	TensorFlow
a = np.zeros((2,2)); b = np.ones((2,2))	a = tf.zeros((2,2)), b = tf.ones((2,2))
np.sum(b, axis=1)	tf.reduce_sum(a, reduction_indices=[1])
a.shape	a.get_shape()
np.reshape(a, (1,4))	tf.reshape(a, (1,4))
b * 5 + 1	b * 5 + 1
np.dot(a,b)	tf.matmul(a, b)
a[0,0], a[:,0], a[0,:]	a[0,0], a[:,0], a[0,:]

DATA TYPES

Data type	Python type	Description
DT_FLOAT	tf.float32	32 bits floating point
DT_DOUBLE	tf.float64	64 bits floating point
DT_INT8	tf.int8	8 bits signed integer
DT_INT16	tf.int16	16 bits signed integer
DT_INT32	tf.int32	32 bits signed integer
DT_INT64	tf.int64	64 bits signed integer
DT_UINT8	tf.uint8	8 bits unsigned integer
DT_STRING	tf.string	Variable length byte arrays. Each element of a Tensor is a byte array
DT_BOOL	tf.bool	Boolean
DT_COMPLEX64	tf.complex64	Complex number made of two 32 bits floating points: real and imaginary parts
DT_QINT8	tf.qint8	8 bits signed integer used in quantized Ops
DT_QINT32	tf.qint32	32 bits signed integer used in quantized Ops
DT_QUINT8	tf.quint8	8 bits unsigned integer used in quantized Ops

CONSTANTS

Constants, Sequences, and Random Values

Constant Value Tensors

```
tf.zeros(shape, dtype=tf.float32,  
name=None)  
  
tf.zeros_like(tensor, dtype=None,  
name=None)  
  
tf.ones(shape, dtype=tf.float32,  
name=None)  
  
tf.ones_like(tensor, dtype=None,  
name=None)  
  
tf.fill(dims, value, name=None)  
  
tf.constant(value, dtype=None,  
shape=None, name='Const')
```

Sequences

```
tf.linspace(start, stop, num,  
name=None)  
  
tf.range(start, limit=None,  
delta=1, name='range')
```

Random Tensors

Examples:

```
tf.random_normal(shape, mean=0.0,  
stddev=1.0, dtype=tf.float32,  
seed=None, name=None)  
  
tf.truncated_normal(shape,  
mean=0.0, stddev=1.0,  
dtype=tf.float32, seed=None,  
name=None)  
  
tf.random_uniform(shape, minval=0,  
maxval=None, dtype=tf.float32,  
seed=None, name=None)  
  
tf.random_shuffle(value,  
seed=None, name=None)  
  
tf.random_crop(value, size,  
seed=None, name=None)  
  
tf.multinomial(logits,  
num_samples, seed=None, name=None)  
  
tf.set_random_seed(seed)
```

Etc.

TENSOR OPERATIONS

Reduction

```
tf.reduce_sum(input_tensor,  
reduction_indices=None,  
keep_dims=False, name=None)  
  
tf.reduce_prod(input_tensor,  
reduction_indices=None,  
keep_dims=False, name=None)  
  
tf.reduce_min(input_tensor,  
reduction_indices=None,  
keep_dims=False, name=None)  
  
tf.reduce_max(input_tensor,  
reduction_indices=None,  
keep_dims=False, name=None)  
  
tf.reduce_mean(input_tensor,  
reduction_indices=None,  
keep_dims=False, name=None)  
  
tf.reduce_all(input_tensor,  
reduction_indices=None,  
keep_dims=False, name=None)  
  
tf.reduce_any(input_tensor,  
reduction_indices=None,  
keep_dims=False, name=None)
```

Basic Math Functions

```
tf.add_n(inputs, name=None)  
tf.abs(x, name=None)  
tf.neg(x, name=None)  
tf.sign(x, name=None)  
tf.inv(x, name=None)  
tf.square(x, name=None)  
tf.round(x, name=None)  
tf.sqrt(x, name=None)  
tf.rsqrt(x, name=None)  
tf.pow(x, y, name=None)
```

Matrix Math Functions

```
tf.trace(x, name=None)  
tf.transpose(a, perm=None,  
name='transpose')  
  
tf.matmul(a, b, transpose_a=False,  
transpose_b=False,  
a_is_sparse=False,  
b_is_sparse=False, name=None)
```

Etc.

TENSORFLOW SESSION

- Session encapsulates environment where TF-objects are evaluated

```
In [8]: import tensorflow as tf
```



```
In [9]: with tf.Session() as sess:  
....:     a = tf.zeros((2,2))  
....:     b = tf.ones((2,2))  
....:     print sess.run(tf.reduce_sum(b, reduction_indices=1))  
....:  
[ 2.  2.]
```

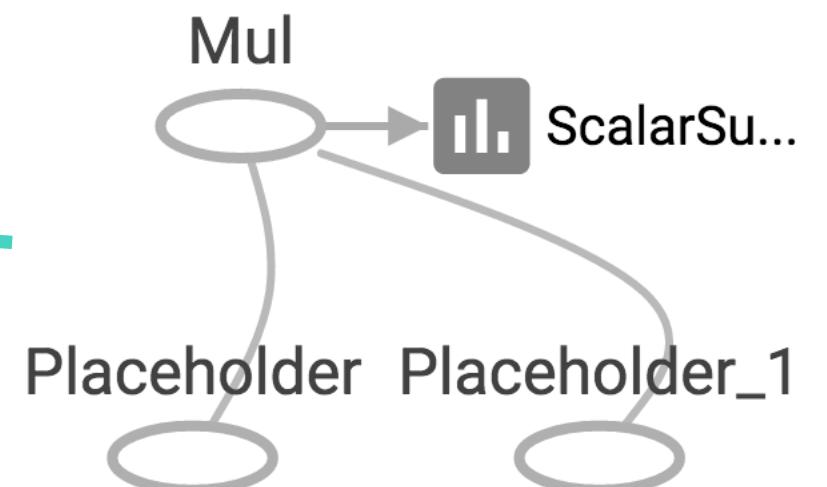


- sess.run() method runs complete subset of graph corresponding to ops passed as arguments

COMPUTATION GRAPH

- Similarly to Theano, we compile a computation graph

```
2 import tensorflow as tf
3
4 with tf.Session() as sess:
5
6     a = tf.placeholder(tf.int32)
7     b = tf.placeholder(tf.int32)
8     c = tf.mul(a, b)
9
10    c_summary = tf.scalar_summary(["c_summary"], c)
11
12    writer = tf.train.SummaryWriter("/tmp/example_logs", sess.graph)
13
14    for i in range(10):
15        feed_dict={a:[2*i], b:[3*i]}
16        summary, c_val = sess.run([c_summary, c], feed_dict=feed_dict)
17        print c_val[0]
18        writer.add_summary(summary, i)
```



TENSORFLOW SESSION - SHORTCUTS

.eval() vs run

```
In [13]: with tf.Session() as sess:  
.....      print tf.ones((1,)).eval()  
.....      print sess.run(tf.ones((1,)))  
.....      print [x for x in sess.run([tf.ones((1,)), tf.ones((1,))])]  
.....  
[ 1.]  
[ 1.]  
[array([ 1.], dtype=float32), array([ 1.], dtype=float32)]
```

!

.eval() can't be used to evaluate multiple values at once

tf.InteractiveSession()

```
In [6]: tf.ones((1,1)).eval()  
-----  
ValueError: Cannot evaluate tensor using eval(): No default session is registered.  
Use `with sess.as_default()` or pass an explicit session to eval(session=sess)
```

```
In [7]: sess = tf.InteractiveSession()  
  
In [8]: tf.ones((1,1)).eval()  
Out[8]: array([[ 1.]], dtype=float32)
```

VARIABLES

Constants

```
In [1]: import tensorflow as tf  
  
In [2]: W1 = tf.ones((2,2))  
  
In [3]: with tf.Session() as sess:  
...:     print W1.eval()  
...:  
[[ 1.  1.]  
 [ 1.  1.]]
```

Variables

```
In [4]: import tensorflow as tf  
  
In [5]: W2 = tf.Variable(tf.ones((2,2)), name="w")  
  
In [6]: with tf.Session() as sess:  
...:     print W2.eval()  
...:  
-----  
-  
FailedPreconditionError
```

Traceback

- "In-memory buffers" containing tensors
- Need to initialize
- Good for weights: they "survive" the execution of the graph

```
In [7]: import tensorflow as tf  
  
In [8]: W2 = tf.Variable(tf.ones((2,2)), name="w")  
  
In [9]: with tf.Session() as sess:  
...:     sess.run(tf.initialize_all_variables())  
...:  
[[ 1.  1.]  
 [ 1.  1.]]
```

UPDATING VARIABLE STATE

```
In [1]: import tensorflow as tf
```

```
In [2]: counter = tf.Variable(0, name = 'counter')
```

```
In [3]: next_value = tf.add(counter, 1)
```



```
In [4]: update = tf.assign(counter, next_value)
```



```
In [5]:
```

```
In [5]: with tf.Session() as sess:  
....:     sess.run(tf.initialize_all_variables())  
....:     print counter.eval()  
....:     update.eval()  
....:     print counter.eval()  
....:
```

```
0
```

```
1
```

EXTERNAL DATA TO TENSORFLOW

From Numpy

```
In [20]: import numpy as np  
  
In [21]: a = np.ones((2,2))  
  
In [22]: tf_a = tf.convert_to_tensor(a)  
  
In [23]: with tf.Session() as sess:  
.....:     print tf_a.eval()  
.....:  
[[ 1.  1.]  
 [ 1.  1.]]
```

From a feed-dictionary

```
In [36]: a = tf.placeholder(tf.int32)  
  
In [37]: b = tf.placeholder(tf.int32)  
  
In [38]: c = tf.mul(a, b)  
  
In [39]: with tf.Session() as sess:  
.....:     print sess.run([c], feed_dict={a:[2, 7], b:[3, 7]})  
.....:  
[array([ 6, 49], dtype=int32)]
```



- Better option for feeding large datasets
- Placeholders as keys, data as values

VARIABLE SCOPING AND REUSE

```
In [1]: import tensorflow as tf
In [2]: sess = tf.Session()
In [3]: w = tf.Variable(0.0, name="w")
In [4]:
In [4]: with tf.variable_scope('a') as scope:
...:     w1 = tf.get_variable("w", [1])
...:
In [5]: with tf.variable_scope('a') as scope:
...:     tf.get_variable_scope().reuse_variables()
...:     w2 = tf.get_variable("w", [1])
...:
In [6]: w1==w2
Out[6]: True
In [7]: w1.name
Out[7]: u'a/w:0'
```

- Scoping helps with naming

```
In [6]: with tf.variable_scope('a') as scope:
...:     w3 = tf.get_variable("w", [1])
ValueError: Variable a/w already exists, disallowed. Did you mean to set reuse=True in VarScope? Originally defined at:
```

- Reuse needs to be explicit

```
In [7]: with tf.variable_scope('b') as scope:
...:     tf.get_variable_scope().reuse_variables()
...:     w4 = tf.get_variable("w1", [1])
ValueError: Variable b/w1 does not exist, disallowed. Did you mean to set reuse=None in VarScope?
```

- Scoping adds prefix to variables names

FORWARD COMPUTATION LINEAR REGRESSION EXAMPLE

```
1 ...
2 A linear regression learning algorithm example using TensorFlow library.
3 Modified from Aymeric Damien's github.com/aymericdamien/TensorFlow-Examples/
4 ...
5
6 import tensorflow as tf
7 import numpy as np
8
9 # Model
10 X = tf.placeholder("float")
11 Y = tf.placeholder("float")
12
13 W = tf.Variable(np.random.randn(), name="weight")
14 b = tf.Variable(np.random.randn(), name="bias")
15
16 pred = tf.add(tf.mul(X, W), b)
17
```

OPTIMIZERS

LINEAR REGRESSION EXAMPLE

```
18 # Training Data
19 train_X = np.asarray([3.3,4.4,5.5,6.71,6.93,4.168,9.779,6.182,7.59,2.167,
20                 7.042,10.791,5.313,7.997,5.654,9.27,3.1])
21 train_Y = np.asarray([1.7,2.76,2.09,3.19,1.694,1.573,3.366,2.596,2.53,1.221,
22                 2.827,3.465,1.65,2.904,2.42,2.94,1.3])
23 n_samples = train_X.shape[0]
24
25 # Parameters
26 learning_rate = 0.01
27 training_epochs = 1000
28 display_step = 50
29
30 # Optimizer
31 cost = tf.reduce_sum(tf.pow(pred-Y, 2))/(2*n_samples)
32 optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
33
34
35 # Launch the graph
36 with tf.Session() as sess:
37
38     sess.run(tf.initialize_all_variables())
39
40     # Fit all training data
41     for epoch in range(training_epochs):
42         for (x, y) in zip(train_X, train_Y):
43             sess.run(optimizer, feed_dict={X: x, Y: y})
44
45     #Display logs per epoch step
46     if (epoch+1) % display_step == 0:
47         c = sess.run(cost, feed_dict={X: train_X, Y:train_Y})
48         print "Epoch:", '%04d' % (epoch+1), "cost=", "{:.9f}".format(c), \
49             "W=", sess.run(W), "b=", sess.run(b)
50
51     print "Optimization Finished!"
52     training_cost = sess.run(cost, feed_dict={X: train_X, Y: train_Y})
53     print "Training cost=", training_cost, "W=", sess.run(W), "b=", sess.run(b), '\n'
```

```
Epoch: 0050 cost= 0.100161634 W= 0.334835 b= 0.188247
Epoch: 0100 cost= 0.097480297 W= 0.329765 b= 0.224718
Epoch: 0150 cost= 0.095108874 W= 0.324997 b= 0.25902
Epoch: 0200 cost= 0.093011498 W= 0.320512 b= 0.291282
Epoch: 0250 cost= 0.091156594 W= 0.316294 b= 0.321625
Epoch: 0300 cost= 0.089516155 W= 0.312327 b= 0.350164
Epoch: 0350 cost= 0.088065267 W= 0.308596 b= 0.377005
Epoch: 0400 cost= 0.086782187 W= 0.305087 b= 0.402251
Epoch: 0450 cost= 0.085647523 W= 0.301787 b= 0.425993
Epoch: 0500 cost= 0.084644087 W= 0.298682 b= 0.448324
Epoch: 0550 cost= 0.083756670 W= 0.295763 b= 0.469328
Epoch: 0600 cost= 0.082971931 W= 0.293017 b= 0.489082
Epoch: 0650 cost= 0.082278028 W= 0.290434 b= 0.507661
Epoch: 0700 cost= 0.081664391 W= 0.288005 b= 0.525135
Epoch: 0750 cost= 0.081121773 W= 0.285721 b= 0.54157
Epoch: 0800 cost= 0.080641970 W= 0.283572 b= 0.557028
Epoch: 0850 cost= 0.080217734 W= 0.281551 b= 0.571567
Epoch: 0900 cost= 0.079842605 W= 0.27965 b= 0.585241
Epoch: 0950 cost= 0.079510957 W= 0.277863 b= 0.598101
Epoch: 1000 cost= 0.079217732 W= 0.276181 b= 0.610197
Optimization Finished!
Training cost= 0.0792177 W= 0.276181 b= 0.610197
```

OPTIMIZERS

- Initialize optimizer

```
30 # Optimizer
31 cost = tf.reduce_sum(tf.pow(pred-Y, 2))/(2*n_samples)
32 optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
33
```

- Setup derivatives graph

```
42 for (x, y) in zip(train_X, train_Y):
43     sess.run(optimizer, feed_dict={X: x, Y: y})
```

- We run the optimizer for all training data
- Behind the scenes: automatic differentiation (similar to Theano)

EXAMPLES: LOGISTIC REGRESSION

```
10 import tensorflow as tf
11
12 # Import MINST data
13 from tensorflow.examples.tutorials.mnist import input_data
14 mnist = input_data.read_data_sets("/tmp/data/", one_hot=True)
15
16 # Parameters
17 learning_rate = 0.01
18 training_epochs = 25
19 batch_size = 100
20 display_step = 1
21
22 # tf Graph Input
23 x = tf.placeholder(tf.float32, [None, 784]) # mnist data image of shape 28*28=784
24 y = tf.placeholder(tf.float32, [None, 10]) # 0-9 digits recognition => 10 classes
25
26 # Set model weights
27 W = tf.Variable(tf.zeros([784, 10]))
28 b = tf.Variable(tf.zeros([10]))
29
30 # Construct model
31 pred = tf.nn.softmax(tf.matmul(x, W) + b) # Softmax
32
33 # Minimize error using cross entropy
34 cost = tf.reduce_mean(-tf.reduce_sum(y*tf.log(pred), reduction_indices=1))
35 # Gradient Descent
36 optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
37
38 # Initializing the variables
39 init = tf.initialize_all_variables()
40
41 # Launch the graph
42 with tf.Session() as sess:
43     sess.run(init)
44
45 # Training cycle
46 for epoch in range(training_epochs):
47     avg_cost = 0.
48     total_batch = int(mnist.train.num_examples/batch_size)
49
50     # Loop over all batches
51     for i in range(total_batch):
52         batch_xs, batch_ys = mnist.train.next_batch(batch_size)
53
54         # Run optimization op (backprop) and cost op (to get loss value)
55         _, c = sess.run([optimizer, cost], feed_dict={x: batch_xs,
56                                                       y: batch_ys})
57
58         # Compute average loss
59         avg_cost += c / total_batch
60
61     # Display logs per epoch step
62     if (epoch+1) % display_step == 0:
63         print "Epoch:", '%04d' % (epoch+1), "cost=", "{:.9f}".format(avg_cost)
64
65     print "Optimization Finished!"
66
67 # Test model
68 correct_prediction = tf.equal(tf.argmax(pred, 1), tf.argmax(y, 1))
69 # Calculate accuracy
70 accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
71 print "Accuracy:", accuracy.eval({x: mnist.test.images, y: mnist.test.labels})
```

EXAMPLES: MLP

```
33 # Create model
34 def multilayer_perceptron(x, weights, biases):
35     # Hidden layer with RELU activation
36     layer_1 = tf.add(tf.matmul(x, weights['h1']), biases['b1'])
37     layer_1 = tf.nn.relu(layer_1)
38     # Hidden layer with RELU activation
39     layer_2 = tf.add(tf.matmul(layer_1, weights['h2']), biases['b2'])
40     layer_2 = tf.nn.relu(layer_2)
41     # Output layer with linear activation
42     out_layer = tf.matmul(layer_2, weights['out']) + biases['out']
43     return out_layer
44
45 # Store layers weight & bias
46 weights = {
47     'h1': tf.Variable(tf.random_normal([n_input, n_hidden_1])),
48     'h2': tf.Variable(tf.random_normal([n_hidden_1, n_hidden_2])),
49     'out': tf.Variable(tf.random_normal([n_hidden_2, n_classes]))
50 }
51 biases = {
52     'b1': tf.Variable(tf.random_normal([n_hidden_1])),
53     'b2': tf.Variable(tf.random_normal([n_hidden_2])),
54     'out': tf.Variable(tf.random_normal([n_classes]))
55 }
56
57 # Construct model
58 pred = multilayer_perceptron(x, weights, biases)
59
60 # Define loss and optimizer
61 cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(pred, y))
62 optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
63
64 # Initializing the variables
65 init = tf.initialize_all_variables()
66
67 # Launch the graph
68 with tf.Session() as sess:
69     sess.run(init)
70
71     # Training cycle
72     for epoch in range(training_epochs):
73         avg_cost = 0.
74         total_batch = int(mnist.train.num_examples/batch_size)
75         # Loop over all batches
76         for i in range(total_batch):
77             batch_x, batch_y = mnist.train.next_batch(batch_size)
78             # Run optimization op (backprop) and cost op (to get loss value)
79             _, c = sess.run([optimizer, cost], feed_dict={x: batch_x,
80                                                               y: batch_y})
81
82         # Compute average loss
83         avg_cost += c / total_batch
```

EXAMPLES: CONVNET

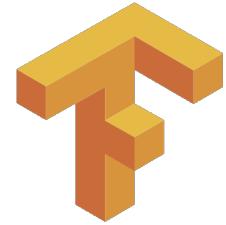
```
34 def conv2d(x, W, b, strides=1):
35     # Conv2D wrapper, with bias and relu activation
36     x = tf.nn.conv2d(x, W, strides=[1, strides, strides, 1], padding='SAME')
37     x = tf.nn.bias_add(x, b)
38     return tf.nn.relu(x)
39
40
41 def maxpool2d(x, k=2):
42     # MaxPool2D wrapper
43     return tf.nn.max_pool(x, ksize=[1, k, k, 1], strides=[1, k, k, 1],
44                           padding='SAME')
45
46
47 # Create model
48 def conv_net(x, weights, biases, dropout):
49     # Reshape input picture
50     x = tf.reshape(x, shape=[-1, 28, 28, 1])
51
52     # Convolution Layer
53     conv1 = conv2d(x, weights['wc1'], biases['bc1'])
54     # Max Pooling (down-sampling)
55     conv1 = maxpool2d(conv1, k=2)
56
57     # Convolution Layer
58     conv2 = conv2d(conv1, weights['wc2'], biases['bc2'])
59     # Max Pooling (down-sampling)
60     conv2 = maxpool2d(conv2, k=2)
61
62     # Fully connected layer
63     # Reshape conv2 output to fit fully connected layer input
64     fc1 = tf.reshape(conv2, [-1, weights['wd1'].get_shape().as_list()[0]])
65     fc1 = tf.add(tf.matmul(fc1, weights['wd1']), biases['bd1'])
66     fc1 = tf.nn.relu(fc1)
67     # Apply Dropout
68     fc1 = tf.nn.dropout(fc1, dropout)
69
70     # Output, class prediction
71     out = tf.add(tf.matmul(fc1, weights['out']), biases['out'])
72     return out
```

EXAMPLES: RNNs

```
49 def RNN(x, weights, biases):
50
51     # Prepare data shape to match `rnn` function requirements
52     # Current data input shape: (batch_size, n_steps, n_input)
53     # Required shape: 'n_steps' tensors list of shape (batch_size, n_input)
54
55     # Permuting batch_size and n_steps
56     x = tf.transpose(x, [1, 0, 2])
57     # Reshaping to (n_steps*batch_size, n_input)
58     x = tf.reshape(x, [-1, n_input])
59     # Split to get a list of 'n_steps' tensors of shape (batch_size, n_input)
60     x = tf.split(0, n_steps, x)
61
62     # Define a lstm cell with tensorflow
63     lstm_cell = rnn_cell.BasicLSTMCell(n_hidden, forget_bias=1.0)
64
65     # Get lstm cell output
66     outputs, states = rnn.rnn(lstm_cell, x, dtype=tf.float32)
67
68     # Linear activation, using rnn inner loop last output
69     return tf.matmul(outputs[-1], weights['out']) + biases['out']
70
71 pred = RNN(x, weights, biases)
```

OTHER FEATURES (NOT COVERED TODAY)

- Checkpoints: can save your model & load it for later use
- Serving: use TF in production
- Details of inner workings of TensorFlow (see academic paper for details)



Background

Usage

Under-the-hood

Resources

OPERATIONS AND KERNELS

OPERATION

Operations have names and represent an abstract computation, such as ‘add’ and ‘matrix multiply’

KERNEL

Implementation of an operation, designed for specific types of devices, such as CPU or GPU

PARALLELIZING

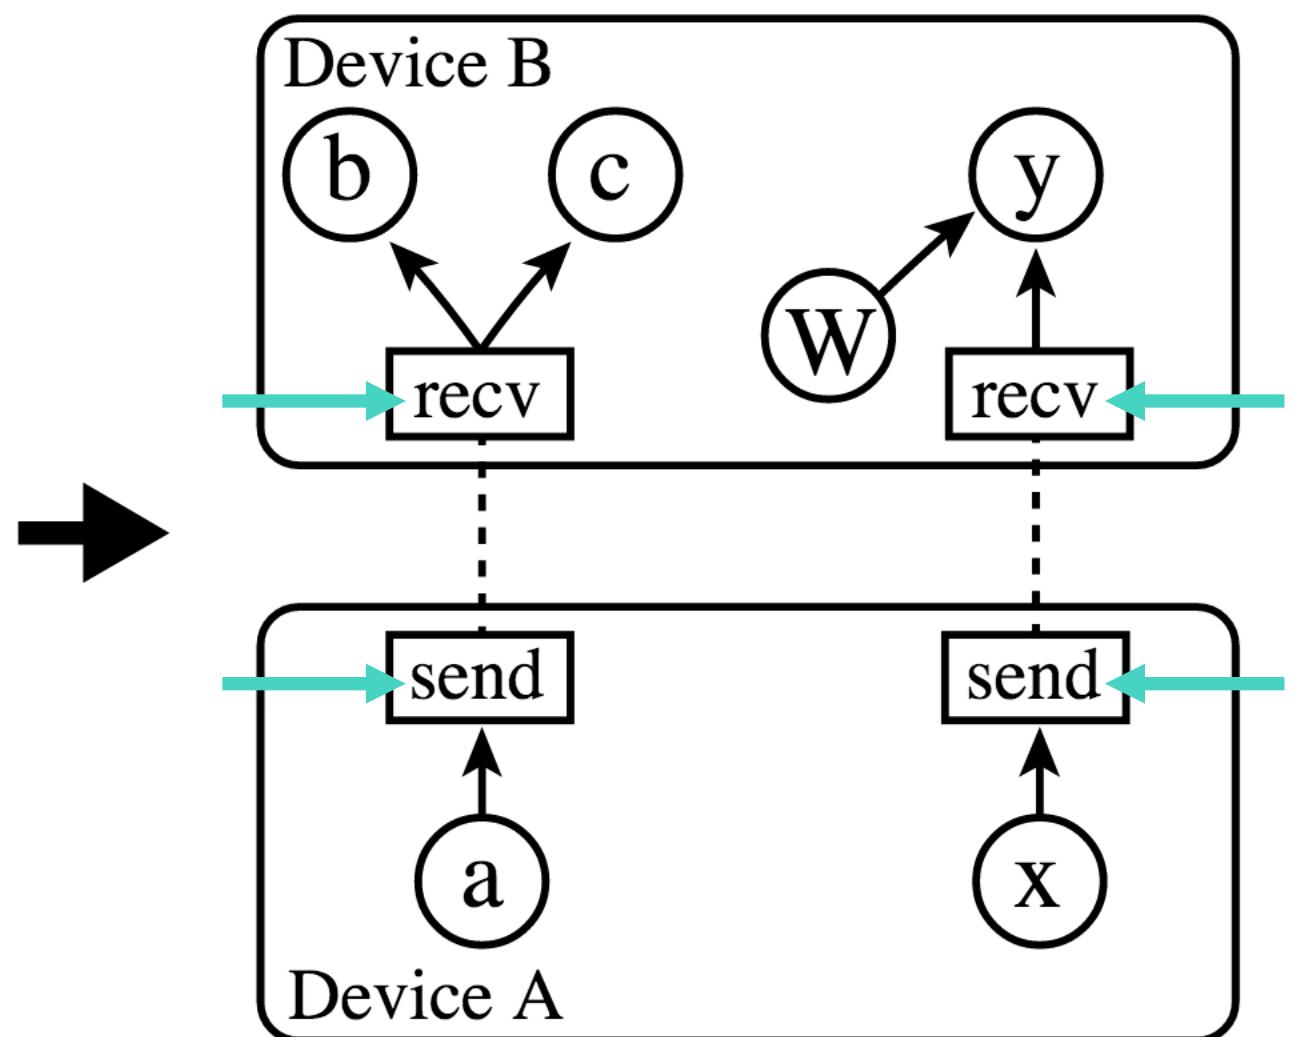
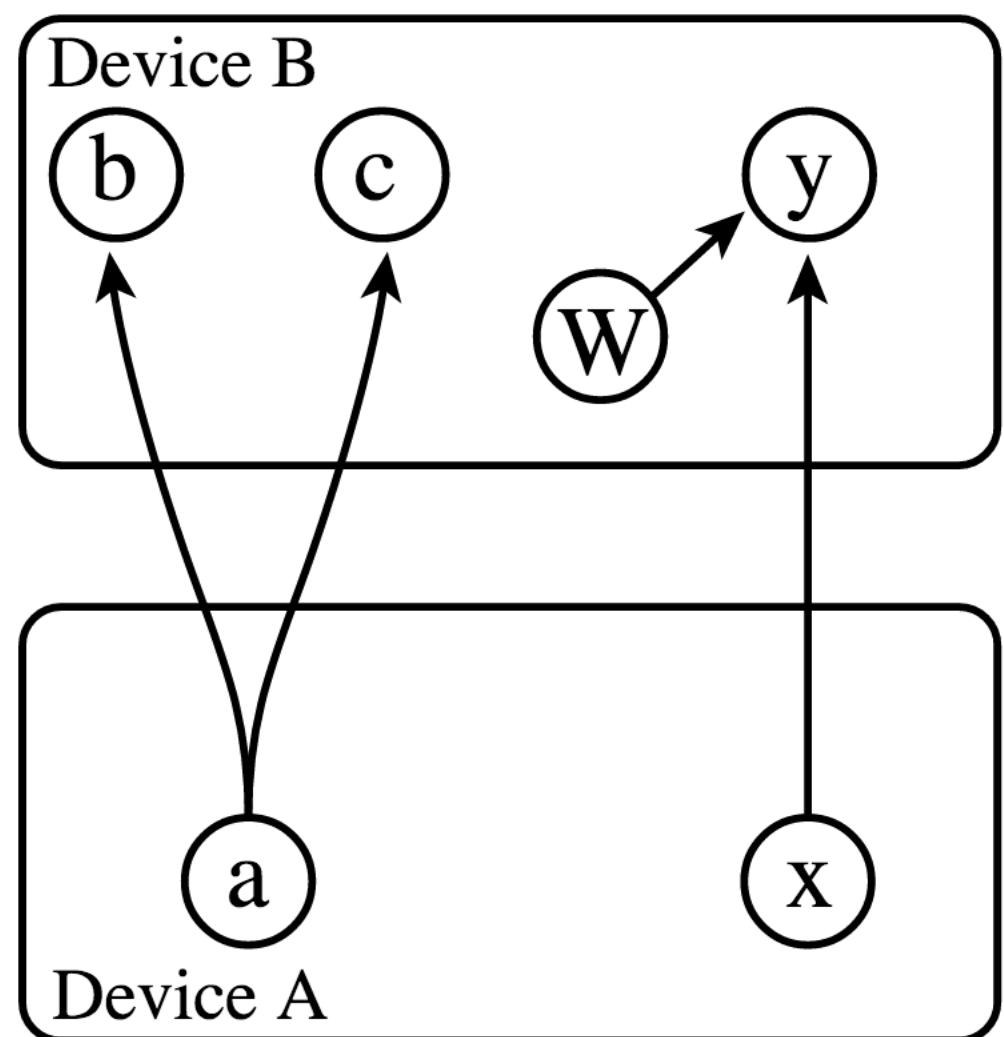
```
In [6]: with tf.device('/cpu:0'):
    ....          a = tf.constant(2)
    ....          b = tf.constant(3)
    ....          c = a*b
    ....
```

```
In [7]: with tf.Session() as sess:  
...:     print c.eval()  
...:  
...:
```

PARALLELIZING: MAPPING ALGORITHM

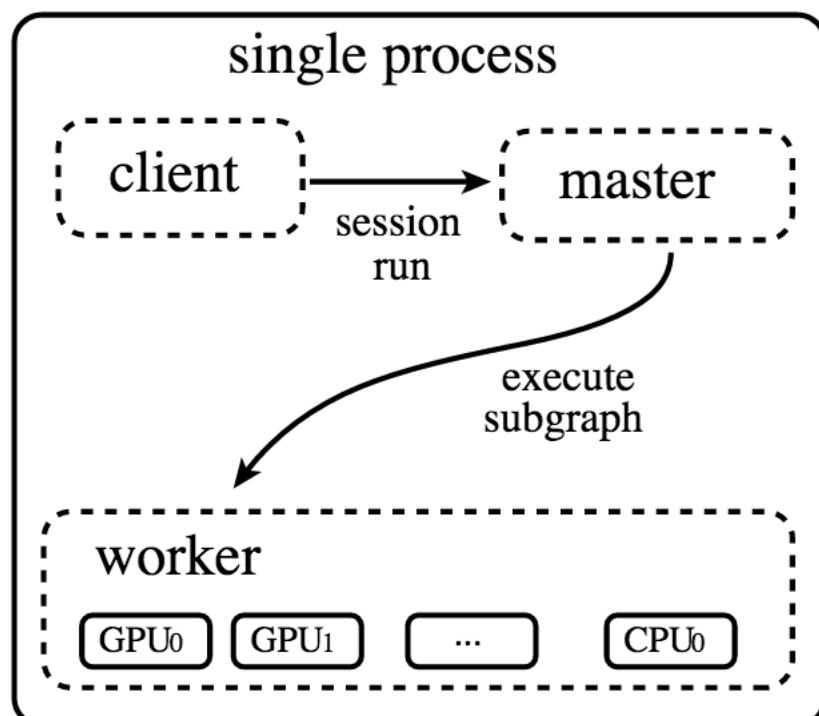
- TF maps computation to available devices
- In summary, during mapping:
 - (1) Algorithm receives cost model
 - (2) Algorithm simulates execution of graph and makes node-placement decisions
 - (3) Real execution is run correspondingly
- Improvements to placement algorithm are in progress

PARALLELIZING: CROSS-DEVICE COMMUNICATION

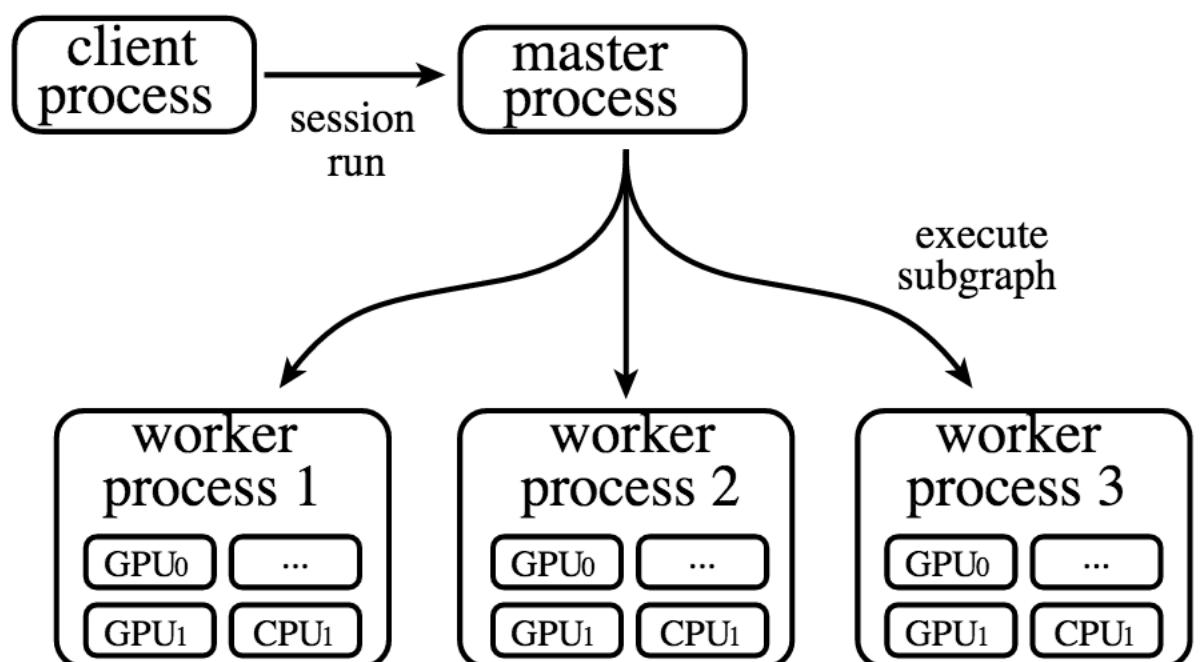


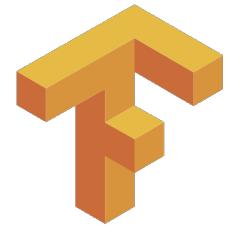
PARALLELIZING: CLIENT, MASTER AND WORKER

1 machine



Distributed system





Background

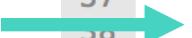
Usage

Under-the-hood

Resources

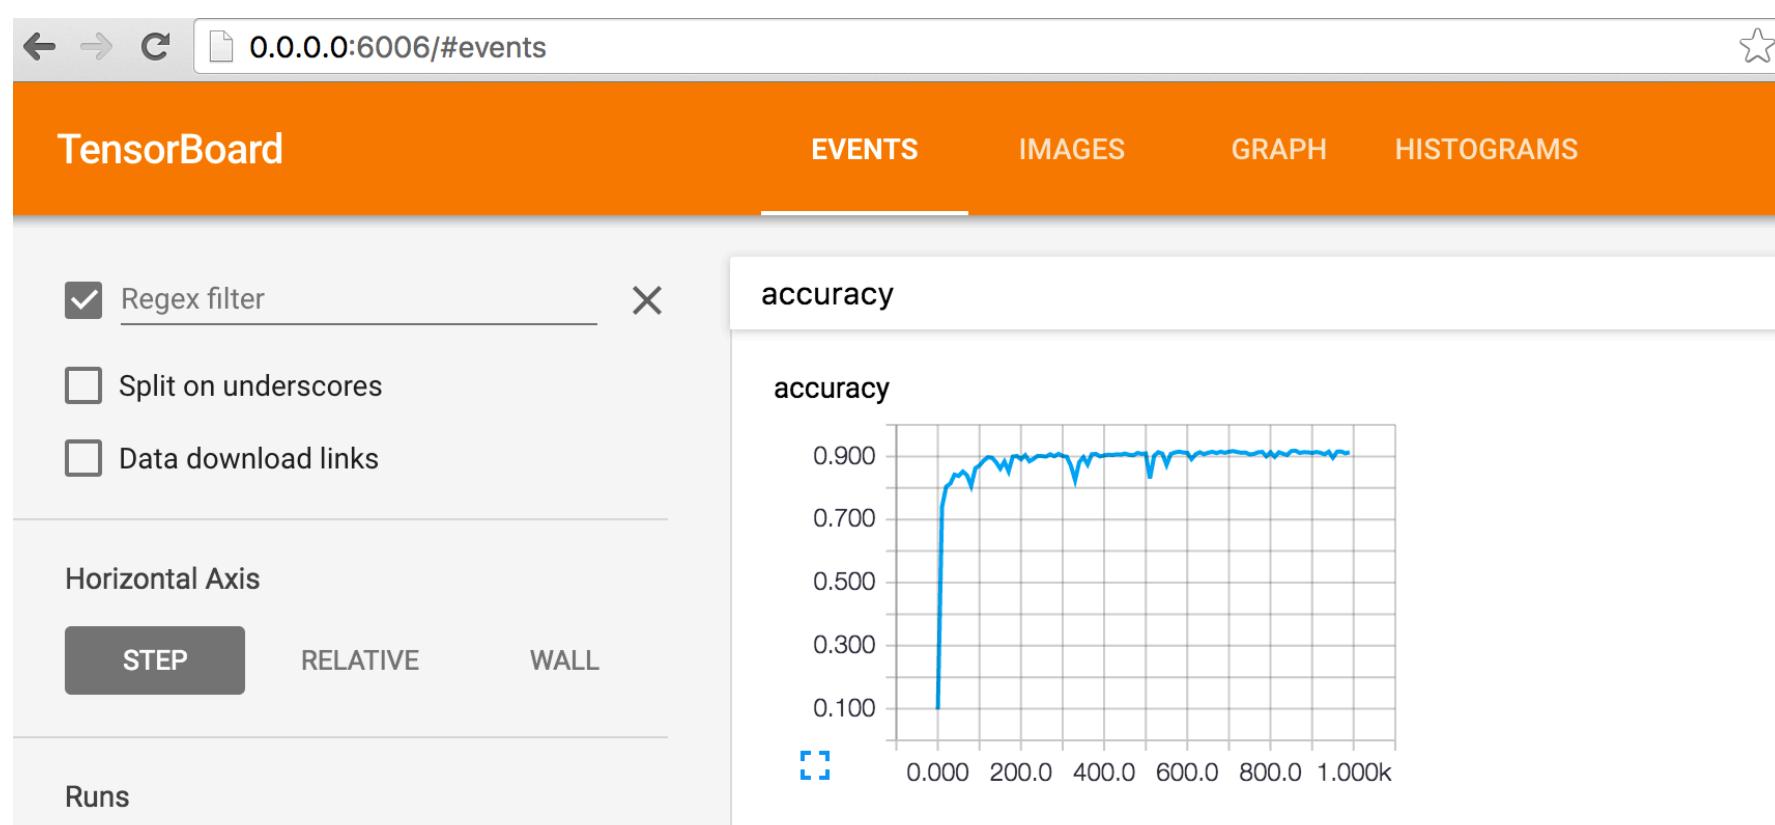
TENSORBOARD: (1) WRITE A SUMMARY

```
2 import tensorflow as tf
3 from tensorflow.examples.tutorials.mnist import input_data
4 mnist = input_data.read_data_sets('MNIST_data', one_hot=True)
5
6 with tf.Session() as sess:
7
8     # Create the model
9     x = tf.placeholder(tf.float32, [None, 784], name="x-input")
10    W = tf.Variable(tf.zeros([784,10]), name="weights")
11    b = tf.Variable(tf.zeros([10]), name="bias")
12
13    # use a name scope to organize nodes in the graph visualizer
14    with tf.name_scope("Wx_b") as scope:
15        y = tf.nn.softmax(tf.matmul(x,W) + b)
16
17    # Add summary ops to collect data
18    w_hist = tf.histogram_summary("weights", W)
19    b_hist = tf.histogram_summary("biases", b)
20    y_hist = tf.histogram_summary("y", y)
21
22    # Define loss and optimizer
23    y_ = tf.placeholder(tf.float32, [None,10], name="y-input")
24    # More name scopes will clean up the graph representation
25    with tf.name_scope("xent") as scope:
26        cross_entropy = -tf.reduce_sum(y_*tf.log(y))
27        ce_summ = tf.scalar_summary("cross entropy", cross_entropy)
28    with tf.name_scope("train") as scope:
29        train_step = tf.train.GradientDescentOptimizer(0.01).minimize(cross_entropy)
30
31    with tf.name_scope("test") as scope:
32        correct_prediction = tf.equal(tf.argmax(y,1), tf.argmax(y_,1))
33        accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
34        accuracy_summary = tf.scalar_summary("accuracy", accuracy)
35
36    # Merge all the summaries and write them out to /tmp/mnist_logs
37    merged = tf.merge_all_summaries()
38    writer = tf.train.SummaryWriter("/tmp/mnist_logs", sess.graph_def)
39    tf.initialize_all_variables().run()
40
41    # Train the model, and feed in test data and record summaries every 10 steps
42
43    for i in range(1000):
44        if i % 10 == 0: # Record summary data, and the accuracy
45            feed = {x: mnist.test.images, y_: mnist.test.labels}
46            result = sess.run([merged, accuracy], feed_dict=feed)
47            summary_str = result[0]
48            acc = result[1]
49            writer.add_summary(summary_str, i)
50            print("Accuracy at step %s: %s" % (i, acc))
51        else:
52            batch_xs, batch_ys = mnist.train.next_batch(100)
53            feed = {x: batch_xs, y_: batch_ys}
54            sess.run(train_step, feed_dict=feed)
55
56    print(accuracy.eval({x: mnist.test.images, y_: mnist.test.labels}))
```



TENSORBOARD: (2) VISUALIZE SUMMARY

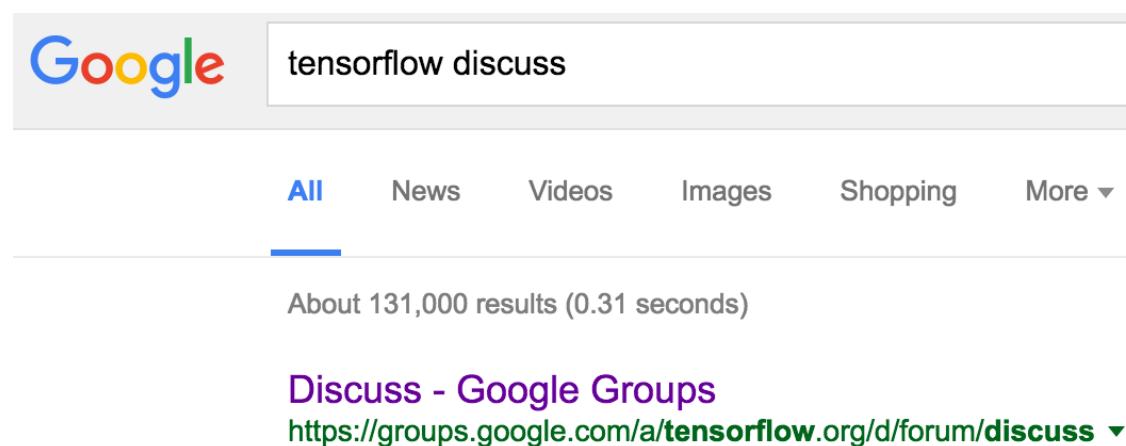
```
MacBook-Pro-60:~ pyry$ tensorboard --logdir=/tmp/mnist_logs
WARNING:tensorflow:Found more than one graph event per run. Overwriting the graph with the newest event.
WARNING:tensorflow:Found more than one graph event per run. Overwriting the graph with the newest event.
Starting TensorBoard 16 on port 6006
(You can navigate to http://0.0.0.0:6006)
```



DEBUGGING TIPS

- Use Numpy: try implementing forward-graph
- Use iPython and tf.InteractiveSession()

- Active community:



- Visualize with TensorBoard
- Review resources (next page)

RESOURCES & REFERENCES

- TF project
 - www.tensorflow.org
 - [www.github.com/tensorflow](https://github.com/tensorflow)
- TF paper + notes
 - ADD PAPER
 - <https://github.com/samjabrahams/tensorflow-white-paper-notes>
- Examples
 - <https://github.com/aymericdamien/TensorFlow-Examples/tree/master/examples>
- Comprehensive list of resources
 - <https://github.com/jtoy/awesome-tensorflow>

HANDS-ON LAB

- [nvlabs.qwiklab.com](https://nvlabs.github.io/qwiklab/tutorials/tensorflow-on-gpus.html) : **Exploring TensorFlow on GPUs**
 - Linear Regression
 - Multilayer Convolutional Network
 - Sequence Autoencoder
 - Multi-GPU Computation
 - Parallelism on GPUs
 - Distributed Training

True AI IS HIRING

SENIOR SOFTWARE ENGINEER

As a Senior Software Engineer, you will be working closely with our researchers and leading the development of technical infrastructure at True AI from ground up, building highly scalable and real time system architecture that blends well with underlying deep learning algorithms

DEEP LEARNING RESEARCH ENGINEER

+ INTERNSHIPS

We need a Deep Learning Research Engineer who is passionate about taking AI to the next level, and who is interested in building the company alongside the founders.

FULL STACK DEVELOPER

We need a Full Stack Developer who is passionate about taking AI to the next level, and who is interested in building the company alongside the founders. You will be playing a key role in the development of our main web application and browser based plugins, as well as integration on backend.

