

# Deep Learning with Torch7

By Abhishek Aggarwal

Co-founder, CTO

True AI

# Table of Contents

1. Introduction to *torch*
  - a. *lua*
  - b. *torch* package
  - c. tensor
2. Building Neural Networks in torch
  - a. *nn*, *optim*
  - b. *nngraph*
3. Writing your custom *nn* module
  - a. Information flow in nn module
  - b. *nn* module API
  - c. Example custom module

Introduction to Torch

Neural networks in Torch

Writing custom neural network module

# Torch7

- A machine learning programming framework built on Lua(JIT), C++, and CUDA
- Open-source, BSD license
- Used by AI labs in industries and academic labs
  - Facebook AI research
  - Google DeepMind (discontinued after release of TensorFlow recently)
  - Twitter Cortex
- Installation
  - **\$ git clone https://github.com/torch/distro.git ~/torch --recursive**
  - **\$ cd ~/torch; bash install-deps; ./install.sh**
- Run Torch7
  - **\$ th main.lua**

# Lua Basics

- Lua is interpreted language like Python and MATLAB
- Learn lua in 15 mins - <http://tylerneylon.com/a/learn-lua/>
- Variables

- `num = 42 -- All numbers are doubles.`
  - `s = 'walternate' -- Immutable strings like Python.`
  - `t = "double-quotes are also fine"`

- Loops

- `while num < 50 do`
  - `num = num + 1`
  - `end`
  - `fredSum = 0`
  - `for j = 100, 1, -1 do fredSum = fredSum + j end`

# Lua Basics (continued)

- Conditional statements

- ```
if num > 40 then
```
- ```
    print('over 40')
```
- ```
elseif s ~= 'walternate' then -- ~= is not equals.
```
- ```
    print('Winter is coming, ' .. line)
```
- ```
end
```
- ```
aBoolValue = false
```
- ```
ans = aBoolValue and 'yes' or 'no' --> 'no'
```

- Functions

- ```
function fib(n)
```
- ```
    if n < 2 then return 1 end
```
- ```
    return fib(n - 2) + fib(n - 1)
```
- ```
end
```

# Lua Basics (continued)

- scope
  - Unless specified with the keyword **local**, every variable or function is global by default
  - **local g; g = function (x) return math.sin(x) end**
- Tables
  - Tables are the only compound data structure, tables are associative arrays
  - Dictionary
  - **t = {key1 = 'value1', key2 = false}**
  - **print(t.key1)** -- Prints 'value1'.
  - **t.newKey = {}** -- Adds a new key/value pair.
  - **t.key2 = nil** -- Removes key2 from the table.
  - Iteration of dictionary
  - **for key, val in pairs(u) do**
  - **print(key, val)**
  - **end**

## Lua Basics (continued)

- Tables (continued)

- `v = {'value1', 'value2', 1.21, 'gigawatts'}`
- `for i = 1, #v do` -- #v is the size of v for lists.
- `print(v[i])` -- Indices start at 1 !! SO CRAZY!
- `end`

- Table as Class

- `Dog = {}`
- `function Dog:new()`
- `newObj = {sound = 'woof'}`
- `self.__index = self`
- `return setmetatable(newObj, self)`
- `end`
- `function Dog:makeSound()`
- `print('I say ' .. self.sound)`
- `end`



# Torch Core Package

- **torch**: tensors, class factory, serialization, BLAS
- **nn**: neural networks, Modules and Criteria
- **optim**: SGD, LBFGS and other optimization functions
- **gnuplot**: plotting and data visualization
- **paths**: functions related paths, and file system
- **image**: save, load, crop, etc.
- **trepl**: the torch LuaJIT interpreter (like lpython)
- **cwrap**: wrapping C/CUDA functions in lua
- Much more at <https://github.com/torch>
- Cheatsheet: <https://github.com/torch/torch7/wiki/Cheatsheet>

# Installation of packages

- Luarocks is the package manager supported by torch

```
$ luarocks install image      # an image library for Torch7
$ luarocks install nnx        # lots of extra neural-net modules
$ luarocks install camera     # a camera interface for Linux/MacOS
$ luarocks install ffmpeg     # a video decoder for most formats
```

# Tensor class

- Tensor is the most important class for math operations.
- Tensor object is a serializable, multidimensional matrix
- Built on top of Storage class, Tensor defines particular way of viewing a storage.
- ```
x = torch.Tensor(5,6)
x:fill(0) -- fill with zeros
x:uniform(-1,1) -- fill with random values between -1 and 1
x:nDimension() -- return 2
x:nElements() -- returns 30
x:size() -- returns torch.LongTensor({5,6})
x:size(2) -- returns 6
y = x[2][4]
```
- Documentation: <https://github.com/torch/torch7/blob/master/doc/tensor.md>
- Documentation: <https://github.com/torch/torch7/blob/master/doc/math.md>

Introduction to Torch

Neural networks in Torch

Writing custom neural network module

## nn package

- Implements most neural network related layers and modules.
- *cunn* is the cuda counterpart of *nn*.
- Documentation: <https://github.com/torch/nn>
- Implemented ~140 kinds of network layer/criterion (April 2016)

```
mlp = nn.Sequential()  
mlp:add( nn.Linear(10, 25) ) -- 10 input, 25 hidden units  
mlp:add( nn.Tanh() ) -- some hyperbolic tangent transfer function  
mlp:add( nn.Linear(25, 1) ) -- 1 output  
  
print(mlp:forward(torch.randn(10)))
```

# Layers and Criteria

- Layers
  - Convolution: `nn.SpatialConvolution`, `nn.SpatialFullConvolution`, etc.
  - Activations: `nn.Sigmoid`, `nn.Tanh`, `nn.Sigmoid`, etc.
  - Normalization: `nn.BatchNormization`, `nn.SpatialBatchNormalization`, etc.
- Criteria (aka loss functions)
  - Classification: `nn.ClassNLLCriterion`, `nn.BCECriterion`
  - Regression: `nn.MSECriterion`, `nn.WeightedMSECriterion`
  - Multi-task: `nn.MultiCriterion`

# Example MLP

- Build MLP

```
require "nn"
mlp = nn.Sequential(); -- make a multi-layer perceptron
inputs = 2; outputs = 1; HUs = 20; -- parameters
mlp:add(nn.Linear(inputs, HUs))
mlp:add(nn.Tanh())
mlp:add(nn.Linear(HUs, outputs))
criterion = nn.MSECriterion()
```

- Forward and Backward pass

```
-- feed it to the neural network and the criterion
criterion:forward(mlp:forward(input), output)
-- train over this example in 3 steps
-- (1) zero the accumulation of the gradients
mlp:zeroGradParameters()
-- (2) accumulate gradients
mlp:backward(input, criterion:backward(mlp.output, output))
-- (3) update parameters with a 0.01 learning rate
mlp:updateParameters(0.01)
```

# Using optim package

```
local optimState = {learningRate=0.01}
require 'optim'
local params, gradParams = model:getParameters()
require 'optim'


for epoch=1,50 do
  local function feval(params)
    gradParams:zero()

    local outputs = model:forward(batchInputs)
    local loss = criterion:forward(outputs, batchLabels)
    local dloss_doutput = criterion:backward(outputs, batchLabels)
    model:backward(batchInputs, dloss_doutput)

    return loss, gradParams
  end
  optim.sgd(feval, params, optimState)
end
```



## nn containers



- Difficult to make arbitrary graphs with containers

# nngraph

- *nn* provides great set of layers and modules for use in DL.
- However, building arbitrary complex networks using containers is difficult.
- *nngraph* is a wrapper around *nn*.
- Documentation: <https://github.com/torch/nngraph>
- Every module can be made node (called *gnode*) in the graph
  - `input = nn.Identity()()`
  - `L1 = nn.Tanh()(nn.Linear(10, 20)(input))`
- Works with all *nn* modules and criterions.
- Allows even printing of forward and backward graph.
- Allows annotation of *gnodes* which helps in debugging.

# nngraph example

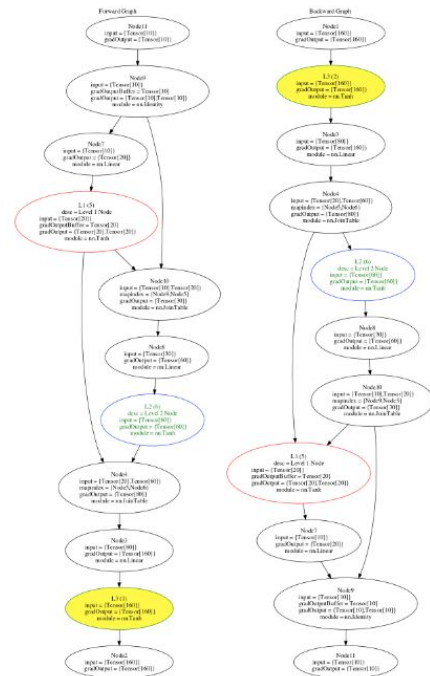
```

input = nn.Identity()()
L1 = nn.Tanh()(nn.Linear(10, 20)(input)):annotate{
  name = 'L1', description = 'Level 1 Node',
  graphAttributes = {color = 'red'}}
}
L2 = nn.Tanh()(nn.Linear(30, 60)(nn.JoinTable(1)({input, L1})):annotate{
  name = 'L2', description = 'Level 2 Node',
  graphAttributes = {color = 'blue', fontcolor = 'green'}}
}
L3 = nn.Tanh()(nn.Linear(80, 160)(nn.JoinTable(1)({L1, L2})):annotate{
  name = 'L3', description = 'Level 3 Node',
  graphAttributes = {color = 'green',
  style = 'filled', fillcolor = 'yellow'}}
}
g = nn.gModule({input},{L3})

indata = torch.rand(10)
gdata = torch.rand(160)
g:forward(indata)
g:backward(indata, gdata)

graph.dot(g.fg, 'Forward Graph', '/tmp/fg')
graph.dot(g.bg, 'Backward Graph', '/tmp/bg')

```



Introduction to Torch


Neural networks in Torch

Writing custom neural network module

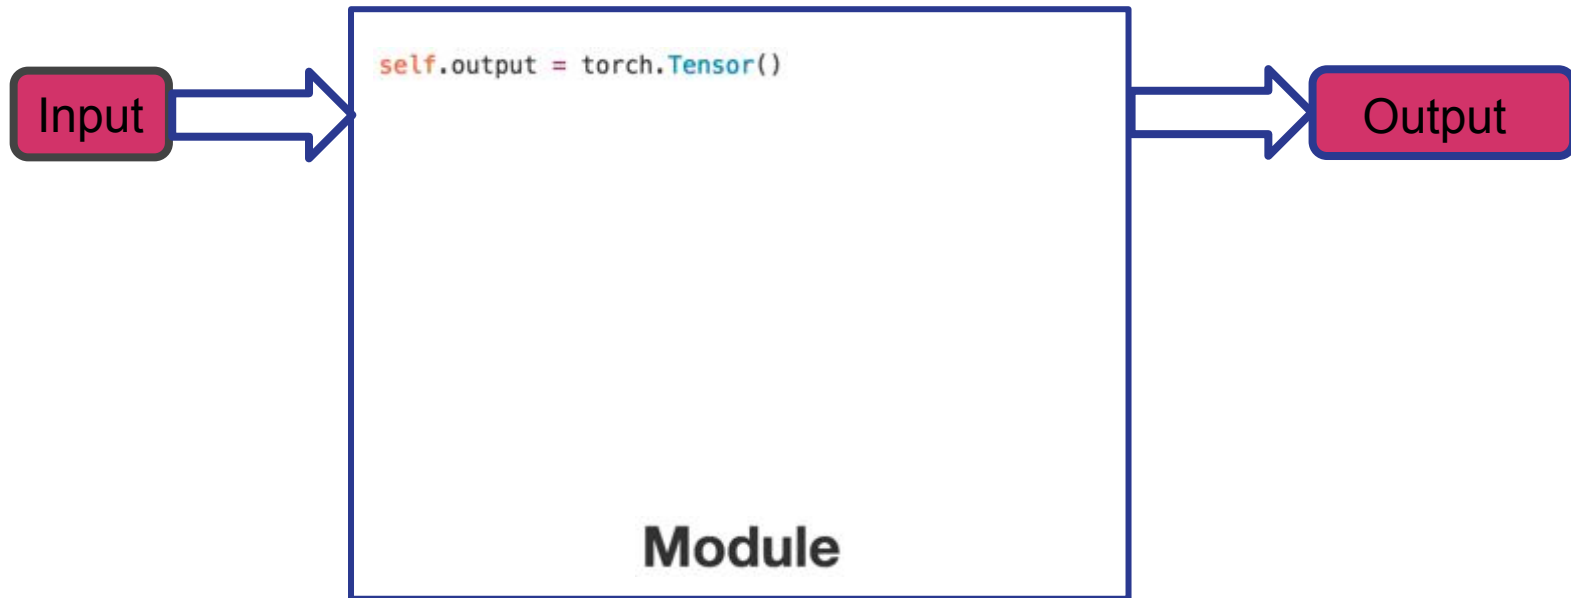
# nn.Module API

- *Module* is an abstract class which defines fundamental methods necessary for training a neural network.
- Everything in nn is inherited from *nn.Module*
- This includes layers, criteria and containers.
- *nn.Module* defines an API required for doing forward and backward pass.
- Thus containers (such *nn.Sequential*) are *nn.Modules* that contain other *nn.Modules* (such as layers and criteria.)
- At high level, all *nn.Module* must implement
  - **forward** - takes input and returns output.
  - **backward** - takes input and updates gradients with respect to input ( and updates gradients of weights, if any)
- Documentation: <https://github.com/torch/nn/blob/master/Module.lua>

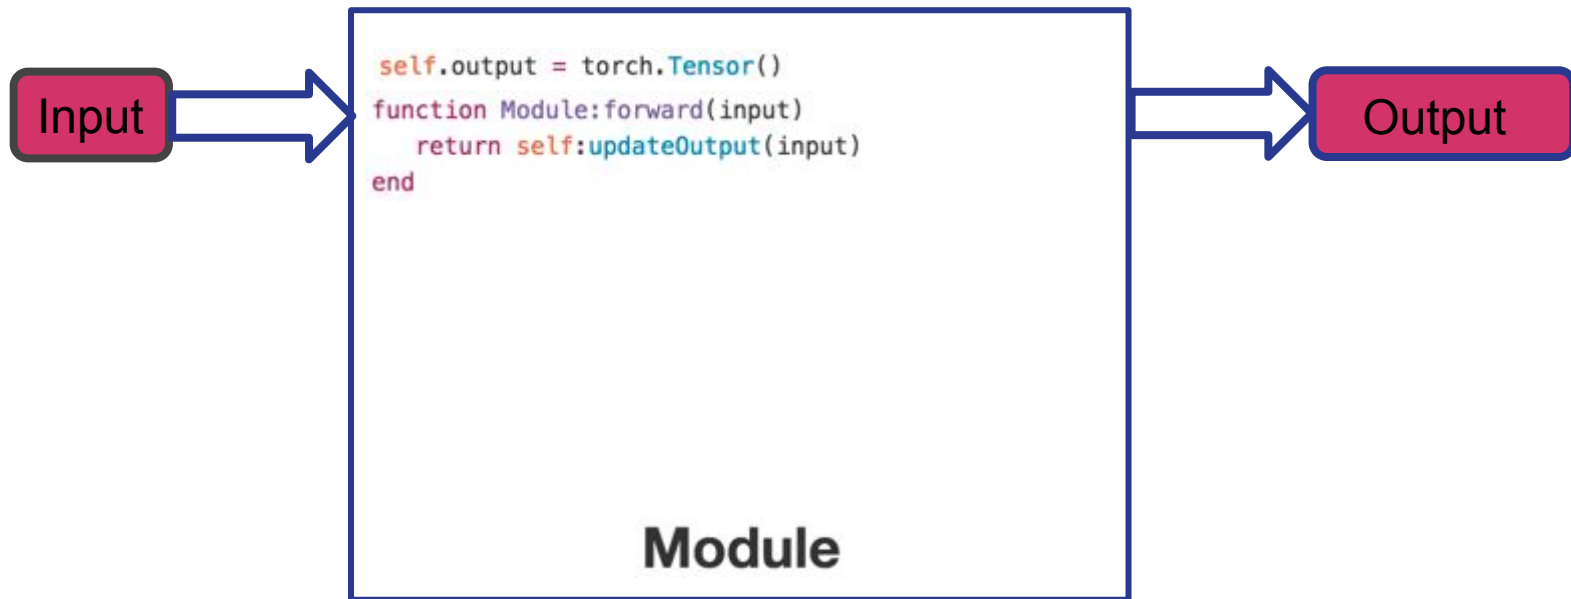
# nn.Module



# nn.Module

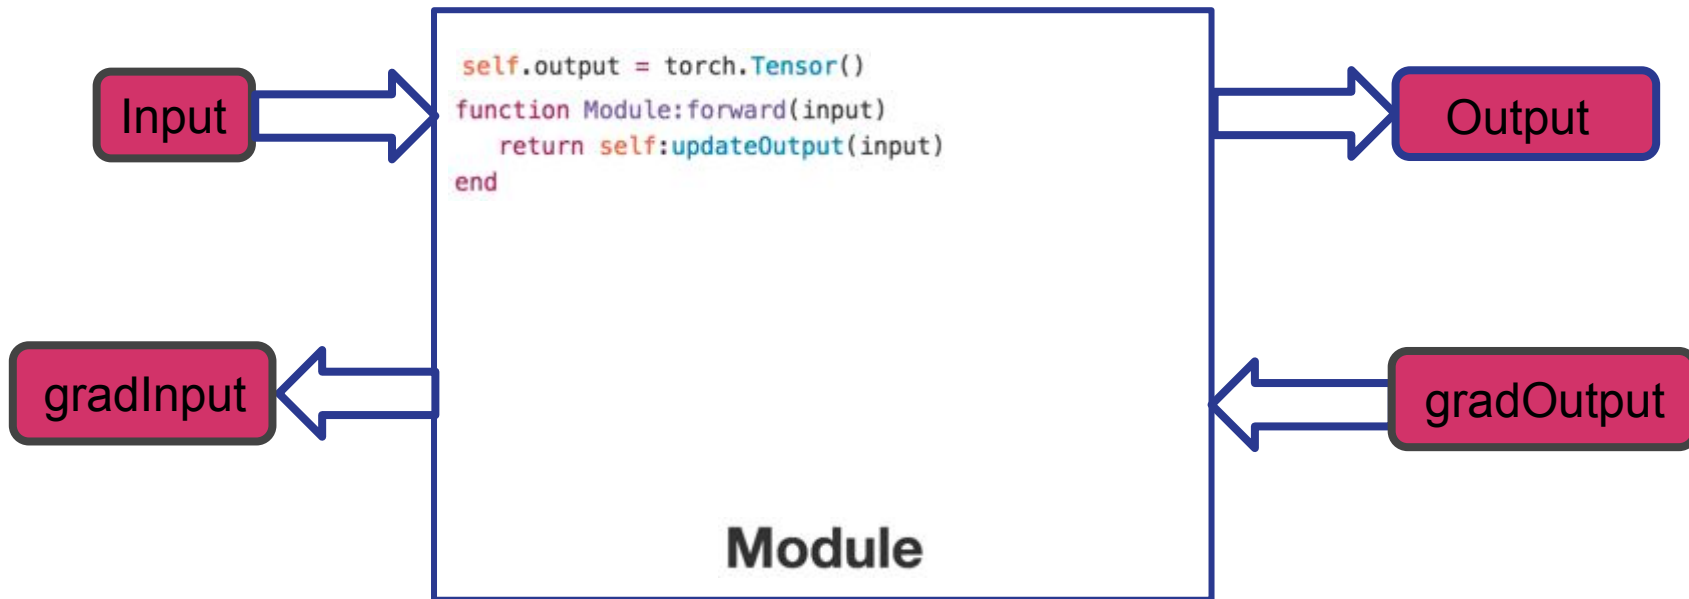


# nn.Module

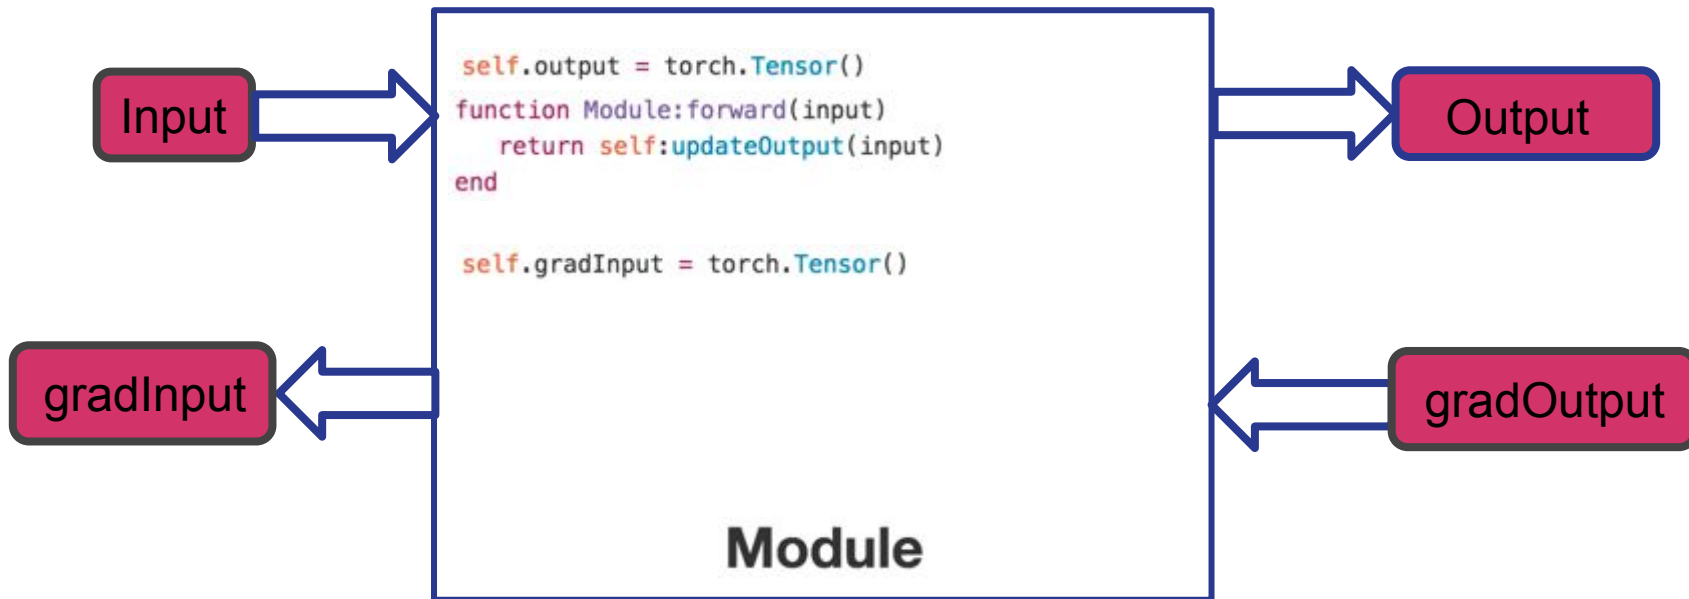




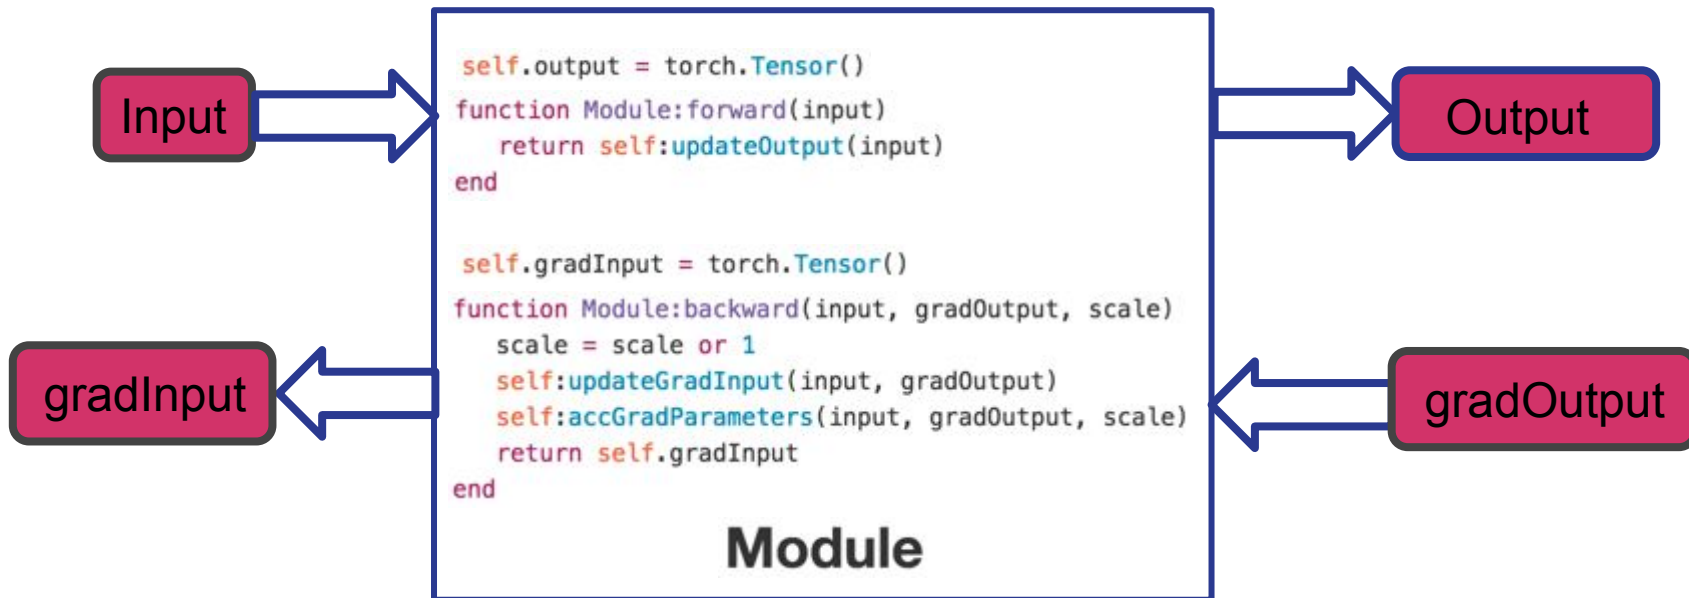
# nn.Module



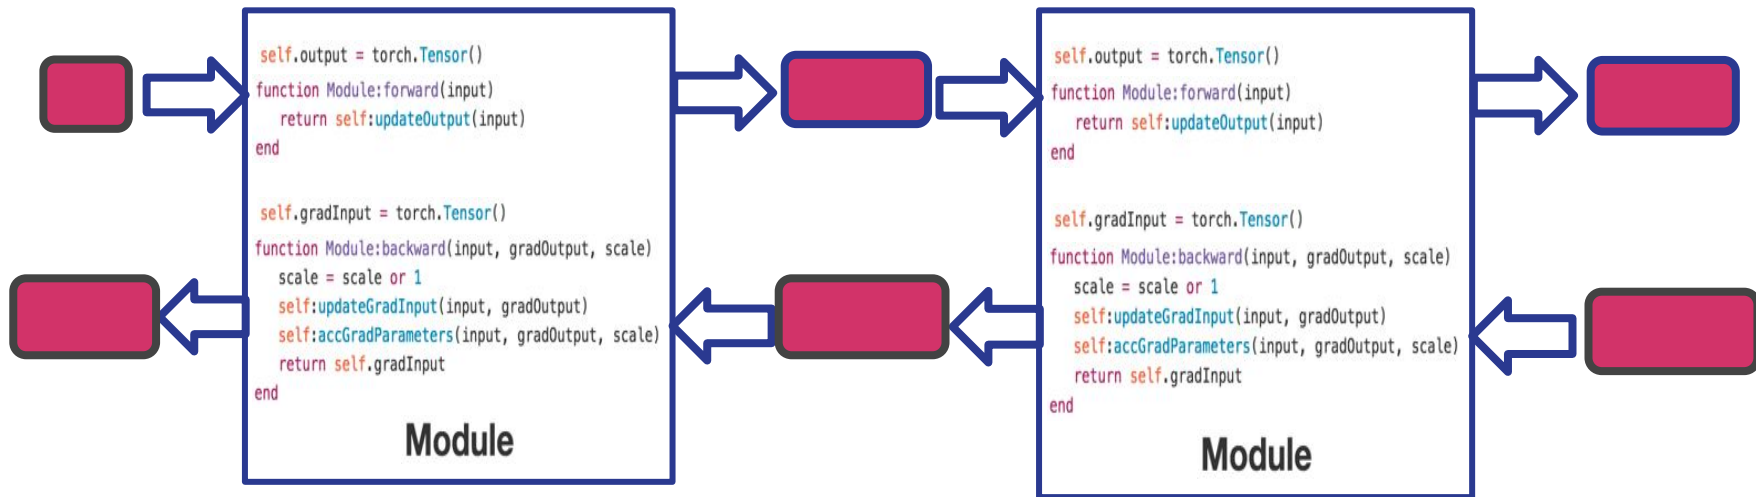
# nn.Module



# nn.Module



# Cascade of nn.Modules



# ReQLu

- We are given  $g_z = \begin{cases} x^2 + x & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$  function, called Rectified Quadratic Linear Unit, or ReQLu.
- Note that this Module has no parameters.
- Implementation:
  - Inherit ReQLu from nn.Module
  - Override updateOutput
  - Override updateGradInput

# ReLU: Forward Pass

```
local ReLU, parent = torch.class('ReLU', 'nn.Module')
-- transfer function f(x) = x^2 + x if x > 0 else 0

function ReLU:__init()
    parent.__init(self)
    -- the two states module needs to maintain are outputs in forward and backward pass
    self.output = torch.Tensor()
    self.gradInput = torch.Tensor()
end

-- define input to output mapping (forward pass)
function ReLU:updateOutput(input)
    -- make sure the input is two dimensional ( batch_size x input_dimension)
    assert(input:nDimension() == 2)

    -- calculate output without mask
    self.output:resizeAs(input):copy(input)
    self.output:cmul(input):add(input)

    -- apply mask
    local mask = input:gt(0):typeAs(input)
    self.output:cmul(mask)
    return self.output
end
```

$$z = \begin{cases} x^2 + x & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

# ReLU: Backward Pass

- Define backward pass

$$\frac{\partial z}{\partial x} = \begin{cases} 2x + 1 & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

$$gradInput = \frac{\partial loss}{\partial x} = \frac{\partial loss}{\partial z} \cdot \frac{\partial z}{\partial x} = gradOutput \cdot \frac{\partial z}{\partial x}$$

```
-- define gradOutput to gradInput mapping (backward pass)
function ReLU:updateGradInput(input, gradOutput)
    self.gradInput:resizeAs(input)

    -- calculate dz/dx (without masking)
    self.gradInput:copy(2*input):add(torch.ones(input:size()))

    -- apply mask
    local mask = input:gt(0):typeAs(input) -- convert from ByteTensor to Tensor
    self.gradInput:cmul(mask)

    -- calculate gradInput by multiplying it with gradOutput
    self.gradInput:cmul(gradOutput)
    return self.gradInput
end
```

# ReLU Scaled

- We are going to create a simple transfer function, called Rectified Quadratic Linear Unit Scaled, or ReQLu Scaled.

$$z = \begin{cases} ax^2 + bx & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

- Two parameters:  $a$  and  $b$
- Implementation:
  - Inherit ReQLu from nn.Module
  - Override *updateOutput*
  - Override *updateGradInput*
  - Override *accGradParameters*



# ReLU Scaled: Implementation

```
local ReLU Scaled, parent = torch.class('ReLU Scaled', 'nn.Module')
-- transfer function f(x) = a*x^2 + b*x if x > 0 else 0

function ReLU Scaled:__init()
    parent.__init(self)
    -- the two states module needs to maintain are outputs in forward and backward pass
    self.output = torch.Tensor()
    self.gradInput = torch.Tensor()

    -- declare weights
    self.a = torch.Tensor(1)
    self.b = torch.Tensor(1)

    -- declare grad of weights
    self.grad_a = torch.Tensor(1)
    self.grad_b = torch.Tensor(1)
end
```

$$z = \begin{cases} ax^2 + bx & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

# ReLU Scaled: Forward Pass

```
-- define input to output mapping (forward pass)
function ReLU Scaled:updateOutput(input)
    -- make sure the input is two dimensional ( batch_size x input_dimension)
    assert(input:nDimension() == 2)

    -- calculate output without mask
    self.output:resizeAs(input):copy(input)
    self.output:cmul(self.a[1] * input)
    self.output:add(self.b[1] * input)

    -- apply mask
    local mask = input:gt(0):typeAs(input)
    self.output:cmul(mask)
    return self.output
end
```

$$z = \begin{cases} ax^2 + bx & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

# ReLU Scaled: Backward Pass I

```
-- define gradOutput to gradInput mapping (backward pass)
function ReLU Scaled:updateGradInput(input, gradOutput)
    self.gradInput:resizeAs(input)

    -- calculate dz/dx (without masking)
    self.gradInput:copy(2*self.a[1] * input):add(self.b[1] * torch.ones(input:size()))

    -- apply mask
    local mask = input:gt(0):typeAs(input) -- convert from ByteTensor to Tensor
    self.gradInput:cmul(mask)

    -- calculate gradInput by multiplying it with gradOutput
    self.gradInput:cmul(gradOutput)
    return self.gradInput
end
```

$$\frac{\partial z}{\partial x} = \begin{cases} 2ax + b & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

$$\text{gradInput} = \frac{\partial \text{loss}}{\partial x} = \frac{\partial \text{loss}}{\partial z} \cdot \frac{\partial z}{\partial x} = \text{gradOutput} \cdot \frac{\partial z}{\partial x}$$

# ReQLuScaled: Backward Pass II

```
function ReQLuScaled:accGradParameters(input, gradOutput)
    -- calculate gradient wrt output
    local grad_a = torch.cmul(input, input)
    local grad_b = input

    -- apply mask
    local mask = input:gt(0):typeAs(input) -- convert from ByteTensor to Tensor
    grad_a:cmul(mask)
    grad_b:cmul(mask)

    -- multiply by gradOutput
    grad_a:cmul(gradOutput)
    grad_b:cmul(gradOutput)

    -- update gradients
    self.grad_a = torch.sum(grad_a)
    self.grad_b = torch.sum(grad_b)
end

-- override the parameters function
function ReQLuScaled:parameters()
    local weights = {self.a, self.b}
    local gradWeights = {self.grad_a, self.grad_b}
    return weights, gradWeights
end
```

$$\frac{\partial z}{\partial a} = \begin{cases} x^2 & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

$$\frac{\partial \text{loss}}{\partial a} = \frac{\partial \text{loss}}{\partial z} \cdot \frac{\partial z}{\partial a} = \text{gradOutput} \cdot \frac{\partial z}{\partial a}$$

$$\frac{\partial z}{\partial b} = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

$$\frac{\partial \text{loss}}{\partial b} = \frac{\partial \text{loss}}{\partial z} \cdot \frac{\partial z}{\partial b} = \text{gradOutput} \cdot \frac{\partial z}{\partial b}$$

# Usage and Code

- The modules that we created can be used just like other modules

```
model = nn.Sequential()  
model.add(nn.Reshape(28*28))  
model.add(nn.Linear(28*28, 225))  
model.add(nn.ReLUScaled())  
model.add(nn.Linear(225, 144))  
model.add(nn.Tanh())  
model.add(nn.Linear(144, 10))  
model.add(nn.LogSoftMax())
```

- Code that uses these modules to classify MNIST digits
  - <https://github.com/abhitopia/TorchTalkDLSummerCampLondon>

## True AI IS HIRING

### SENIOR SOFTWARE ENGINEER

As a Senior Software Engineer, you will be working closely with our researchers and leading the development of technical infrastructure at True AI from ground up, building highly scalable and real time system architecture that blends well with underlying deep learning algorithms

### DEEP LEARNING RESEARCH ENGINEER

We need a Deep Learning Research Engineer who is passionate about taking AI to the next level, and who is interested in building the company alongside the founders.

### FULL STACK DEVELOPER

We need a Full Stack Developer who is passionate about taking AI to the next level, and who is interested in building the company alongside the founders. You will be playing a key role in the development of our main web application and browser based plugins, as well as integration on backend.

+ INTERNSHIPS

